

The Dictionary Abstract Data Type

Goodrich et al. chapter 8

The Dictionary ADT

- Dictionary data structures enable us to store, search for and retrieve elements from a collection based on some key value associated with them.
- We have already studied some data structures which have this ability.
 - Last semester we studied the hash table ADT (recall we stored a student in a hash table based on using the X number as a key)
 - This semester we have studied the associative container `map` in the STL
 - Each of these is an example of a dictionary
- An item to be stored in a dictionary has the form (k, e) where k is the key and e is the element
- Dictionaries may be defined to have unique keys
 - as in the STL `map`
- Or they may be defined to store multiple items with the same key (duplicates)
 - The STL provides a template class called `multimap` which supports duplicate keys

Ordered and Unordered Dictionaries

- The essential feature of a dictionary is that it enables us to look things up based on a key
- Dictionaries may be ordered or unordered
 - In either case we use a key as an identifier that is assigned to each element by a user or application.
- In the ordered dictionary, items are stored in an order which is based on their key values.
 - a ‘total order’ relation must be defined on the keys ...
 - ... means we understand ‘greater than’, ‘less than’ as well as ‘equal to’
 - The dictionary must be able to support operations like
`closestBefore(k)`
`closestAfter(k)`
- In the unordered dictionary, no such ordering exists.
 - we can only test for equality of keys.
 - The dictionary can be implemented using an unordered vector, list or general sequence to store the key-value pairs

Basic operations supported by a Dictionary ADT

- *Note, these are not the names of the methods provided by the `map` and `multimap` STL template classes, but methods with similar functionality exist*
 - `size()`
 - `isEmpty()`
 - `elements()` *returns an iterator to the elements stored in the map*
 - `keys()` *returns an iterator to the keys stored in the map*
 - `find(k)` *return position of item with key k, or null*
 - `findAll(k)` *returns an iterator of positions for all items with key k*
 - `insertItem(k, e)`
 - `removeElement(k)` *an error condition if no such k found in map*
 - `removeAllElements(k)` *removes some no. of elements (may be 0)*
- Notice that none of these methods assume any notion of an ‘order’ in the way the items are stored.
 - They only assume that it is possible to check that the key of an item matches a given key

Dictionaries: performance analysis

- The dictionary implemented with an unordered sequence can **insert elements in constant time**
 - Because no sorting based on the key is needed
- But removal or retrieval requires scanning through the entire sequence
 - **Thus worst case running time for retrieval is $O(n)$**
- The unordered dictionary is a good implementation in cases where there are many insertions into the dictionary, but relatively few extractions
- In an ordered dictionary, elements are stored in an order based on their key value
- The key value pairs may be stored in a vector referred to as a look-up table
- **Insertions may take $O(n)$ time**
 - If we need shift all items along to accommodate it
- Retrievals are very fast, because we can use a binary search
 - **Which is $O(\log n)$**
- The ordered dictionary is a better option if there are relatively fewer insertions than there are retrievals from the dictionary

Unordered Dictionaries - log files

- Unordered dictionaries are also referred to as log files or audit trails
- Consider situations where we wish to archive structured data
 - For example, financial database systems storing all their transactions
 - OS programs such as web servers storing all requests they process over the internet
- Such archive or logging requirements are often implemented with a dictionary realised by using an *un-ordered* vector, list or general sequence to store key-element pairs
- And such an implementation is often called a log file or audit trail
- Archiving database and operating system transactions have precisely the requirements provided by an unordered dictionary
 - Insertions must be simple and fast
 - Retrievals are based on a key value
 - But are only needed when something goes wrong, so performance is not such an issue

Ordered Dictionaries: improving insertion and retrieval times

- The ordered dictionary can never equal the log-file for insertions
 - That would be constant time
- Retrievals are very fast, because we can use a binary search
 - Which is $O(\log n)$
- But insertions may take $O(n)$ time if we use an ordered vector to store the data
- Ideally, we need a different data structure than a vector to use within the dictionary implementation
 - which will improve insertion time ...
 - ... without compromising the retrieval time
- Insertion time can be improved if the data is stored internally in a binary search tree instead of a vector
 - If the BST is well balanced Insertions and retrievals may both be done in $O(\log n)$ time, a huge improvement
 - But it may not remain well-balanced, It depends on the order in which items are inserted
 - Insertions and retrievals from a BST in the worst case are $O(n)$
- Our next topic will explore other tree structures to use for a dictionary – for example a balanced binary search tree